

CS 487/587 Database Implementation

Winter 2021

Database Benchmarking Project - Part 2

Team members: **Deep Patel** **Goutham Podili**

In the experiment, we are evaluating PostgreSQL with different memory options and query planner configuration in Postgres server on Google Cloud VM and local computer.

SYSTEM CONFIGURATION:

The memory level configuration of the system is mentioned below.

Machine type	RAM	OS Image	Postgres version
Google VMs	2 GB	Ubuntu-2004	Psq1 12.6
Local	16GB	Windows10	Psq1 13.1

SYSTEM RESEARCH:

Available postgres memory options and query planner.

<code>shared_buffers(integer)</code>	<ul style="list-style-type: none">• Sets the amount of memory the database server uses for shared memory buffers.• Default is typically 128 MB.
<code>work_mem(integer)</code>	<ul style="list-style-type: none">• Specifies the amount of memory to be used by internal sort operations and hash tables before writing to temporary disk files.• Default is 4MB.
<code>temp_buffers(integer)</code>	<ul style="list-style-type: none">• Sets the maximum number of temporary buffers used by each database session.• Default is 8MB.
<code>enable_seqscan(boolean)</code>	<ul style="list-style-type: none">• It is impossible to suppress sequential scans entirely, but turning this variable

	off discourages the planner from using one if there are other methods available. <ul style="list-style-type: none"> • Default is on
<code>enable_hashjoin(boolean)</code>	<ul style="list-style-type: none"> • Enables or disables the query planner's use of hash-join plan types. • Default is on.

PERFORMANCE EXPERIMENTS:

1) Testing the 10% rule of thumb

- This test explores when it is good to use an unclustered index vs not using an index vs using a clustered index. We will also vary the selectivity of queries like 5%, 25%, and 50%.
- We are using relations with size of 10,000 tuples(TENKTUP1) and 100,000 tuples(HUNDREDKTUP).
- We are using Wisconsin benchmark queries 2, 4 and 6
 - INSERT INTO TMP SELECT * FROM TENKTUP1 WHERE unique2 BETWEEN 792 AND 1791 (**no index**)
 - INSERT INTO TMP SELECT * FROM TENKTUP1 WHERE unique2 BETWEEN 792 AND 1791 (**clustered index**)
 - INSERT INTO TMP SELECT * FROM TENKTUP1 WHERE unique1 BETWEEN 792 AND 1791 (**non-clustered index**)
- No parameters are changed in this test.
- We are expecting that for queries which has selectivity less than ten percent attributes which have no index will perform better than attributes with index.

2) Testing work_mem

- This test examines the relationship between the sort operations and work_mem. We will set work_mem to different sizes.
- We are using relations with size of 10,000 tuples(TENKTUP1) and 100,000 tuples(HUNDREDKTUP).
- The queries we are going to execute for this experiment are
 - SELECT unique1, strin3 FROM TENKTUP1 ORDER BY string3.
 - INSERT INTO TMP SELECT * FROM ONEKTUP, TENKTUP1 WHERE (ONEKTUP.unique1 = TENKTUP1.unique1) AND (TENKTUP1.unique1 = TENKTUP2.unique1) AND (TENKTUP1.unique1 < 1000)
- We will change the work_mem parameter in this test.
- As the sorting uses work memory we are expecting to see an increase in the performance of the query when increasing the work_mem.

3) Testing shared_buffer parameter

- a) This test illustrates what size of the shared buffer performs well on given queries.
We will be testing with 25 MB, 128 MB, and 256 MB.
- b) We consider table size of 10,000 tuples(TENKTUP1), 100,000 tuples(HUNDREDKTUP), and 1,000,000 tuples(ONEMILLIONTUPLES)
- c) Queries:
 - i) Join query
SELECT onemilliontuples.unique1 , hundredktuples.unique1
FROM onemilliontuples, hundredktuples
WHERE onemilliontuples.string3 = hundredktuple.string3;
 - ii) Aggregate query
SELECT COUNT(DISTINCT unique1)
FROM TENKTUP1
- d) We will be testing on 25 MB, 128 MB, and 256 MB. Total RAM size of Postgres server is 2 GB
- e) It is recommended to set 25% of RAM for buffer size.
When performing join and aggregate queries with different memory sizes, relatively low memory will provide better results.

4) Testing temp_buffers parameter

- a) This test demonstrates what size of temp buffers is well suited for sort and hash table joins operations queries. We will be testing 4 MB, 8 MB, and 16 MB.
- b) We consider table size of 10,000 tuples(TENKTUP1), 100,000 tuples(HUNDREDKTUP), and 1,000,000 tuples(ONEMILLIONTUPLES)
- c) Queries:
 - i) SELECT TENKTUP1.unique1, HUNDREDKTUP.unique1
FROM TENKTUP, HUNDREDKTUP.
WHERE TENKTUP1.string3 = HUNDREDKTUP..string3;
 - ii) SELECT HUNDREDKTUP.unique1 , ONEMILLIONTUPLES.unique1
FROM HUNDREDKTUP, ONEMILLIONTUPLES
WHERE ONEMILLIONTUPLES.string3 = HUNDREDKTUP.string3;
- d) The experiment will be held on 4 MB, 8 MB, and 16 MB temp_buffer size
- e) Having a bigger temp_buffer size help for sort and hash table operations.

LESSON LEARNED:

- 1) Unable to set shared_buffers size on running psql using command

```
set shared_buffers="256";  
Error: parameter "shared_buffers" cannot be changed without  
restarting the server
```

Since shared buffers is heart of data base system, it only possible to set up from config file and need to restart the server after changing the size.

REFERENCES

<https://www.postgresql.org/docs/9.6/runtime-config-resource.html>